



Advances in Machine Learning at Microsoft Threat Protection

5/29/2019

Christian Seifert, Principal Researcher



Security Research Superheroes

Security Research
Experts



+

Machine Learning
Systems



=

Security Research
Superheroes





Microsoft Defender ATP

Built-in. Cloud-powered.

PRE-BREACH



SMARTSCREEN

Protect against malicious URLs and downloads



ENDPOINT PROTECTION

Protect against all types of emerging threats

POST-BREACH



ENDPOINT DETECTION & RESPONSE

Detect, investigate, and respond to advanced attacks

GOAL

Block at first sight

FP TOLERANCE

Low



SCOPE



www, files, "fileless"

GOAL

Alert on all possible breaches

FP TOLERANCE

Moderate



SCOPE

Same as pre-breach + cross-service + bad actor behaviors

Intelligent Security Graph

Microsoft internal signals

Office 365 ATP Bing
File detonation
Azure ATP
Windows Store

External signals

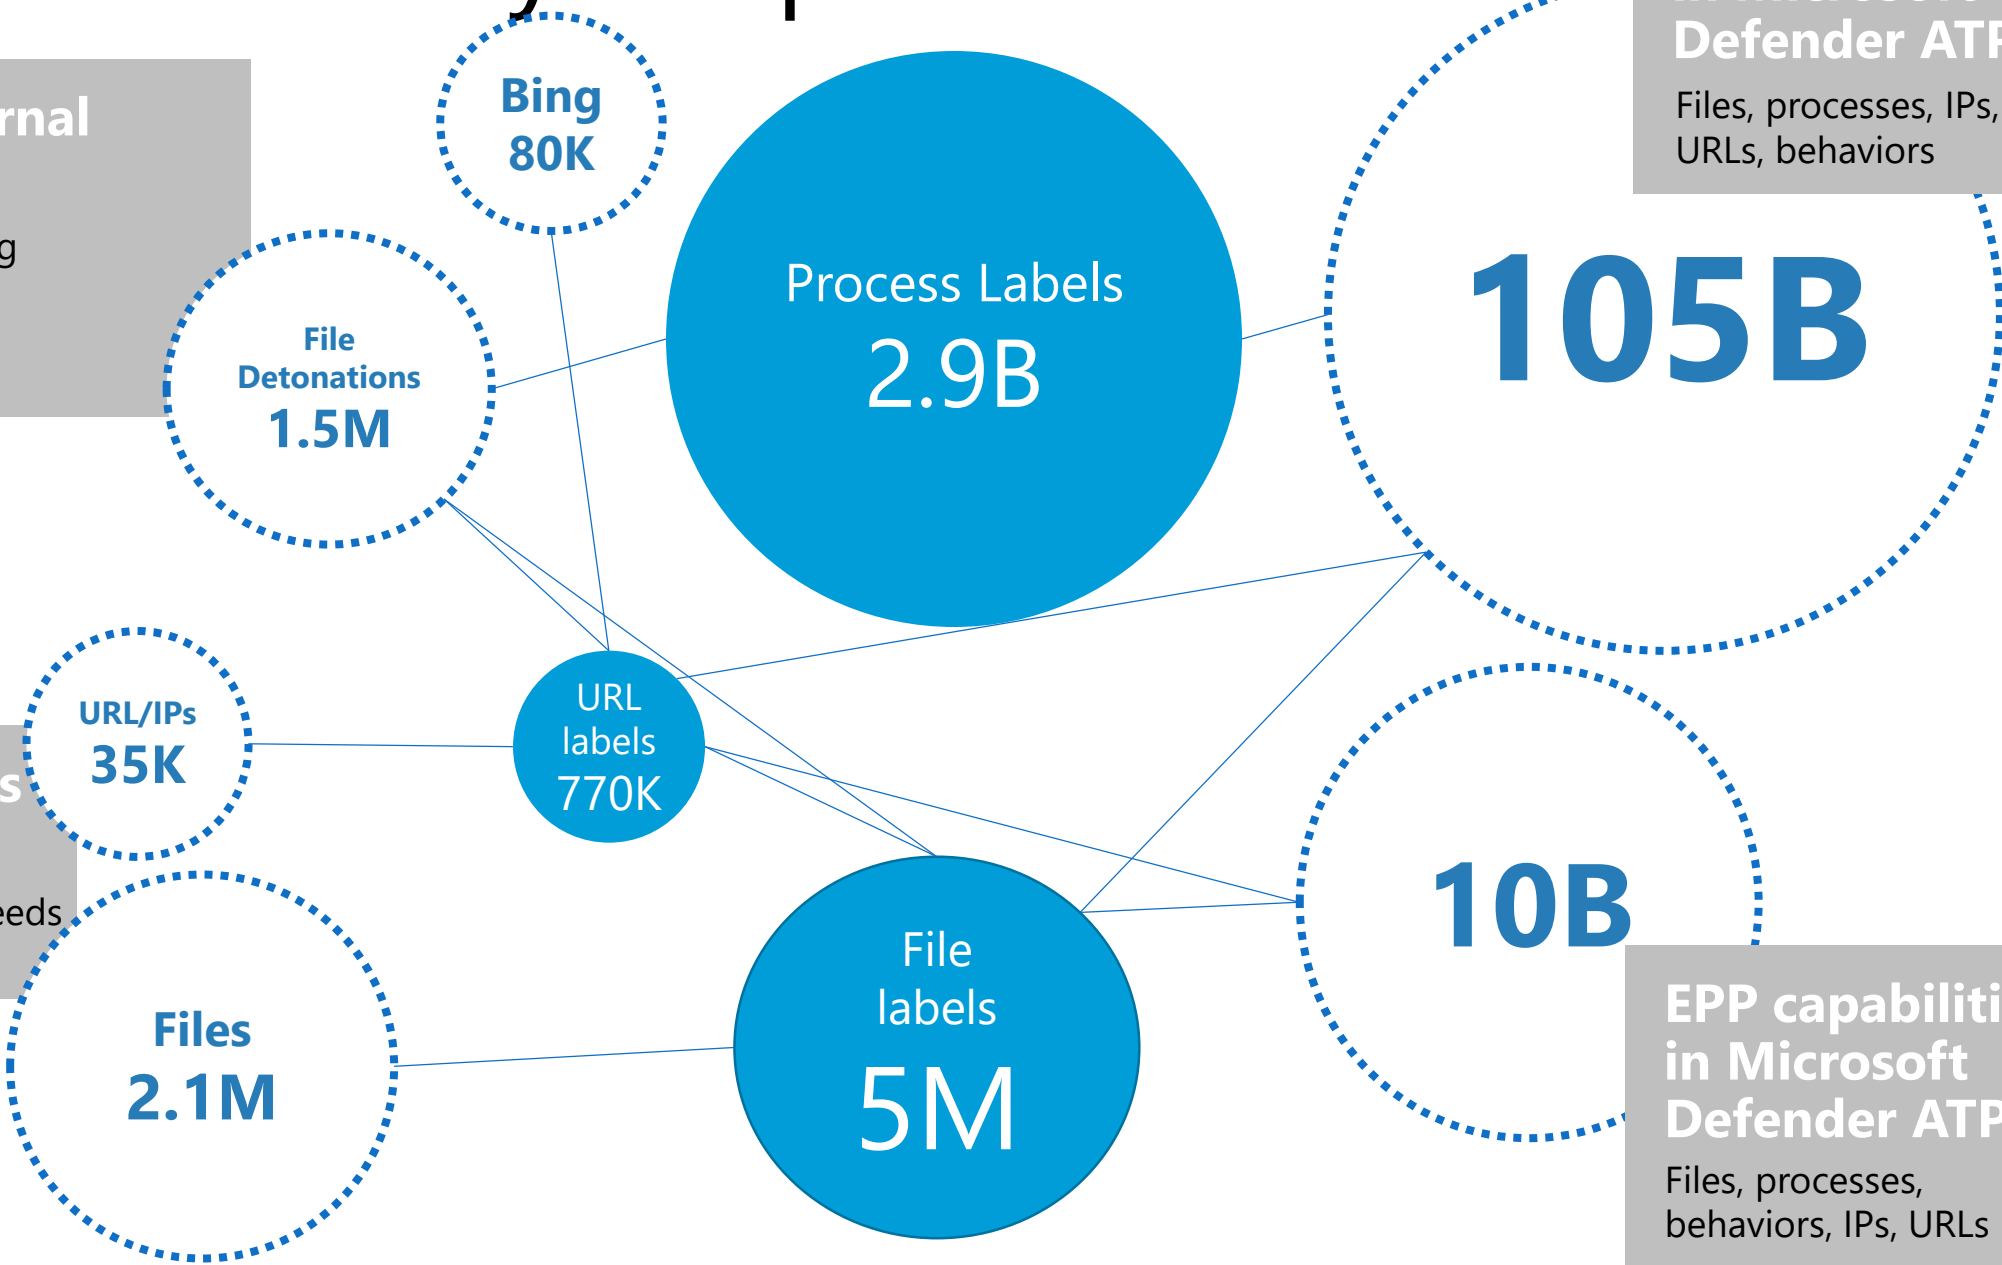
File feeds
URL, domain, and IP feeds

EDR capabilities in Microsoft Defender ATP


Files, processes, IPs, URLs, behaviors

EPP capabilities in Microsoft Defender ATP

Files, processes, behaviors, IPs, URLs



AI diversity



Supervised learning

Fast learners

Deep learning



Unsupervised learning

Anomaly detection


Embeddings



New learners/ approaches

Active learning

Homomorphic
encryption



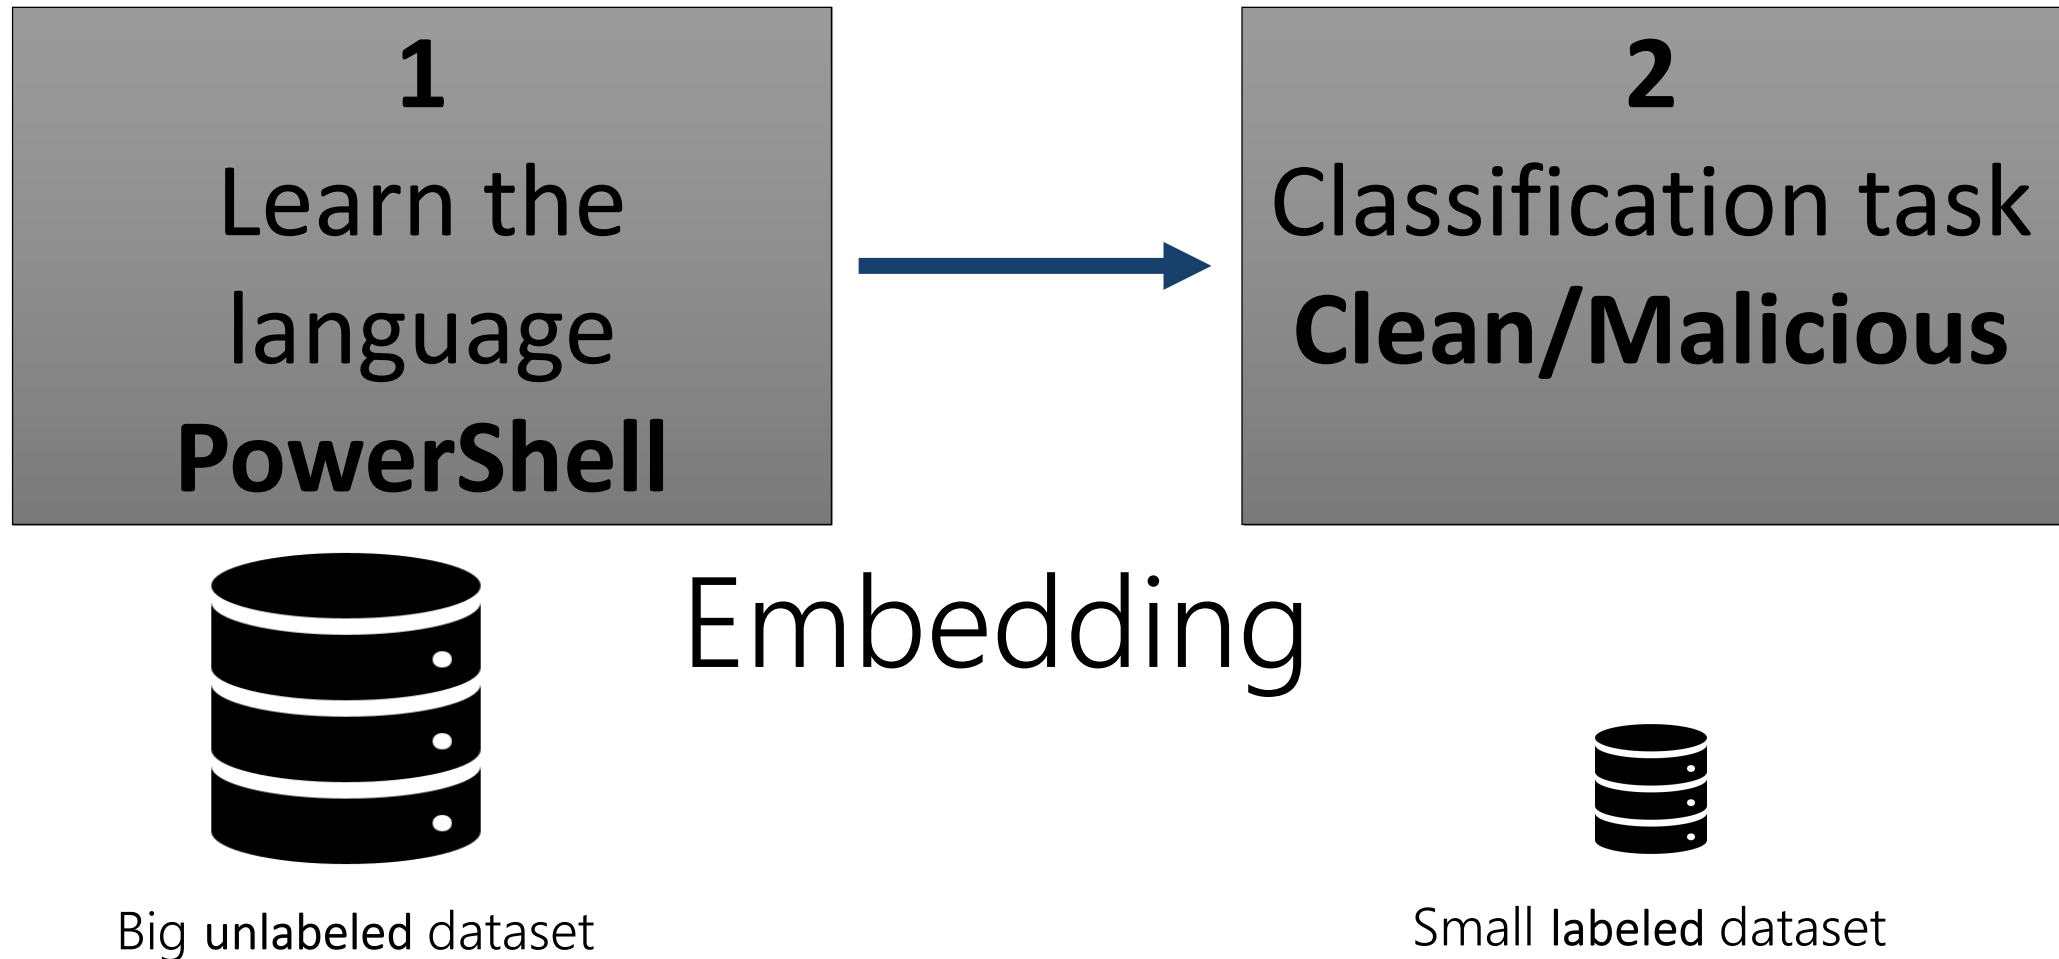
Deep learning in MDATP

Case Study 1: Deep learning for malicious PowerShell detection

Why PowerShell?

Why Deep learning?

Case Study 1: Deep learning for malicious PowerShell detection




Contextual Embedding - **English** (NLP)

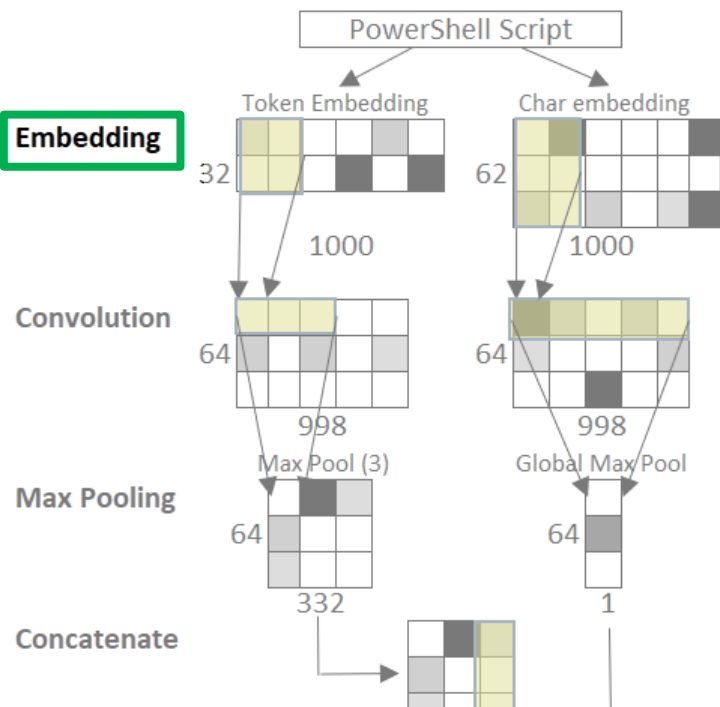
- Tokens \longrightarrow Vectors (\mathbb{R}^n)
- king - man + woman = queen

Contextual Embedding - PowerShell

- High - \$false + \$true = Low
- 'Export-CSV' - \$csv + \$html = 'ConvertTo-Html'
- 'Get-Process' - \$processes + \$services = 'Get-service'



Learning the semantics of PowerShell



True Positive rate:
+22%
improvement

V1 in production
(Using ONNX
and ML.Net)


Paper published



Supervised learners

Case Study 2: Monotonic model

Problem: Adversaries!




malicious features \times clean features

Solution: Monotonic approach only weights malicious features

Incer, Inigo, et al. "Adversarially robust malware detection using monotonic classification." *Proceedings of the Fourth ACM International Workshop on Security and Privacy Analytics*. ACM, 2018.

Over the past month...

0.25M additional blocks



Monotonic model

TECHNOLOGY NEWS MARCH 26, 2019 / 8:20 AM / UPDATED 6 HOURS AGO

Norsk Hydro's initial loss from cyber attack may exceed \$40 million

Nerijus Adomaitis

3 MIN READ



OSLO (Reuters) - Norwegian aluminum maker Norsk Hydro may have lost more than \$40 million in the week that followed a cyber attack that paralyzed parts of its operations, and a full recovery of IT systems will take weeks or more, the company said.

| | |
|---------------------------------|--|
| Sha256 | c97d9bbc80b573bdeeda3812f4d00e5183493dd0d5805e2508728f65977dda15 |
| Determination | Malware |
| Signer | ALISA LTD |
| Age | 0 |
| Prevalence | 0 |
| FilePredictV4_Malware | 34.4% |
| FilePredictV4_Clean | 27.7% |
| FilePredictV4_Malware_Monotonic | 78.8% |
| WinningRuleName | FilePredictV4_Malware_Monotonic |

Monotonic model ignores clean features (certificates with positive rep)

Signature Info ⓘ

Signature Verification

 A certificate was explicitly revoked by its issuer.

File Version Information

| | |
|---------------|------------------------------|
| Copyright | Copyright (C) ALISA LTD 2019 |
| Product | Service tgytutrc |
| Description | Background Tasks Host |
| Original Name | tgytutrc |
| Internal Name | tgytutrc |
| File Version | 1.5.1.0 |
| Date Signed | 2:11 PM 3/21/2019 |

Signers

-  ALISA LTD
-  Sectigo RSA Code Signing CA
-  USERTrust Secure™




Anomaly detection


Statistical anomaly detection



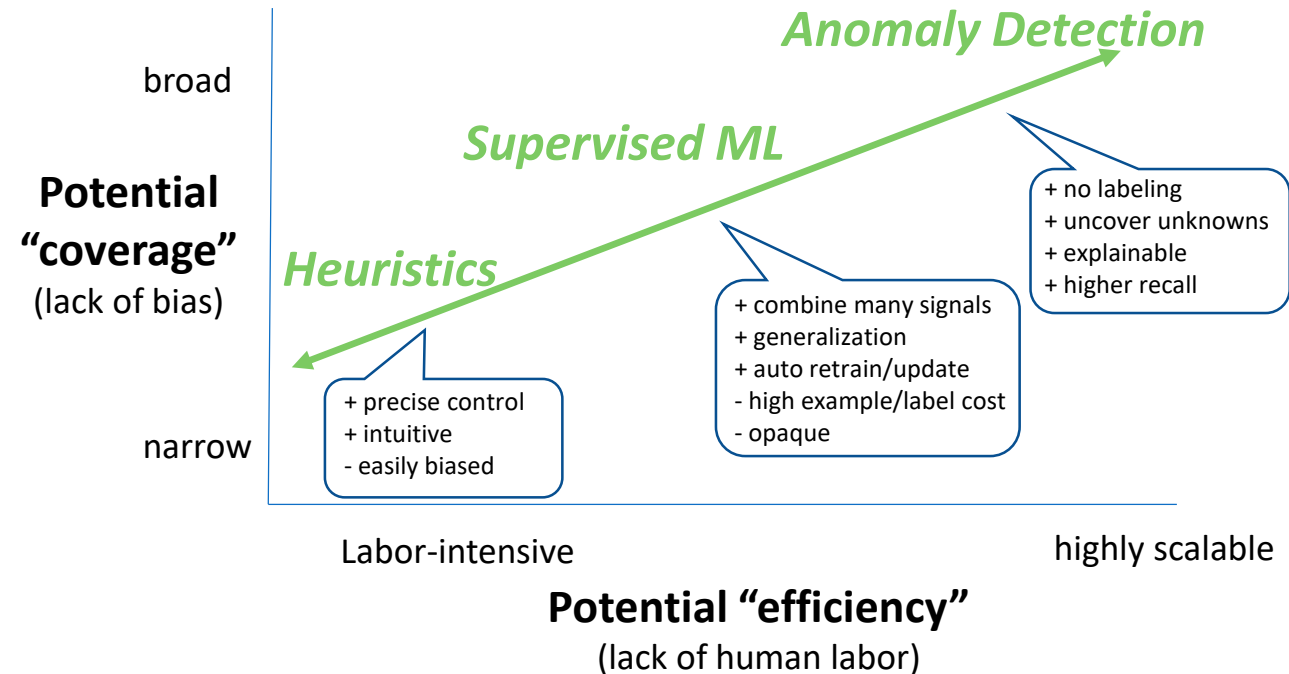
Categorical



Population



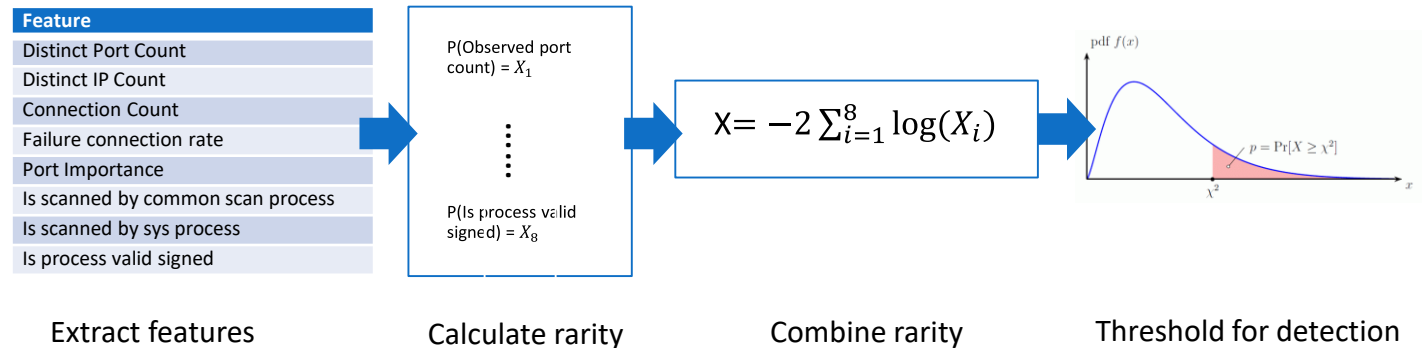
Time-Series



Case Study 3: Port scan anomaly detection

Identify internal attack reconnaissance

- Both vertical and horizontal port scanning detector
- 90%+ precision

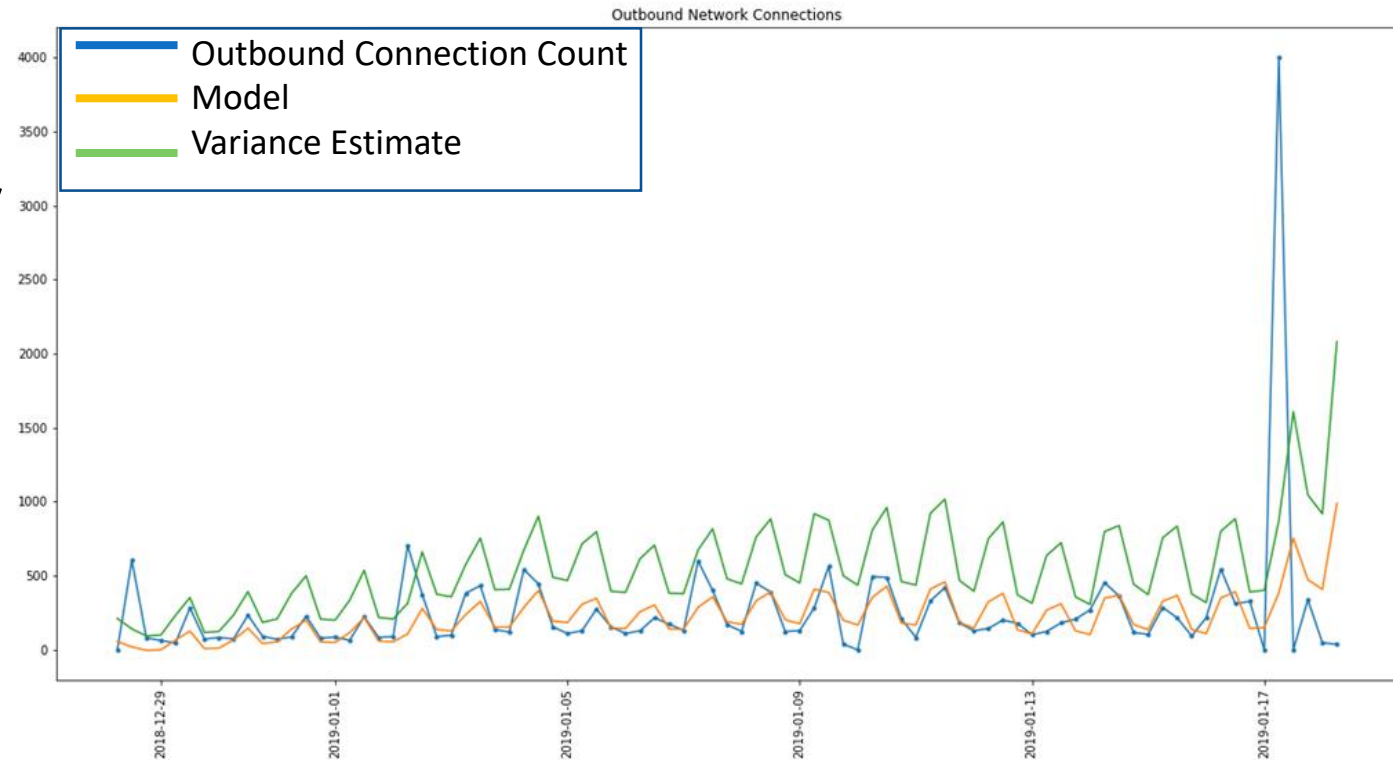



Identified Emotet and Monero ransomware malware.

Case Study 4: Brute force time series anomaly detection

Model of login behavior

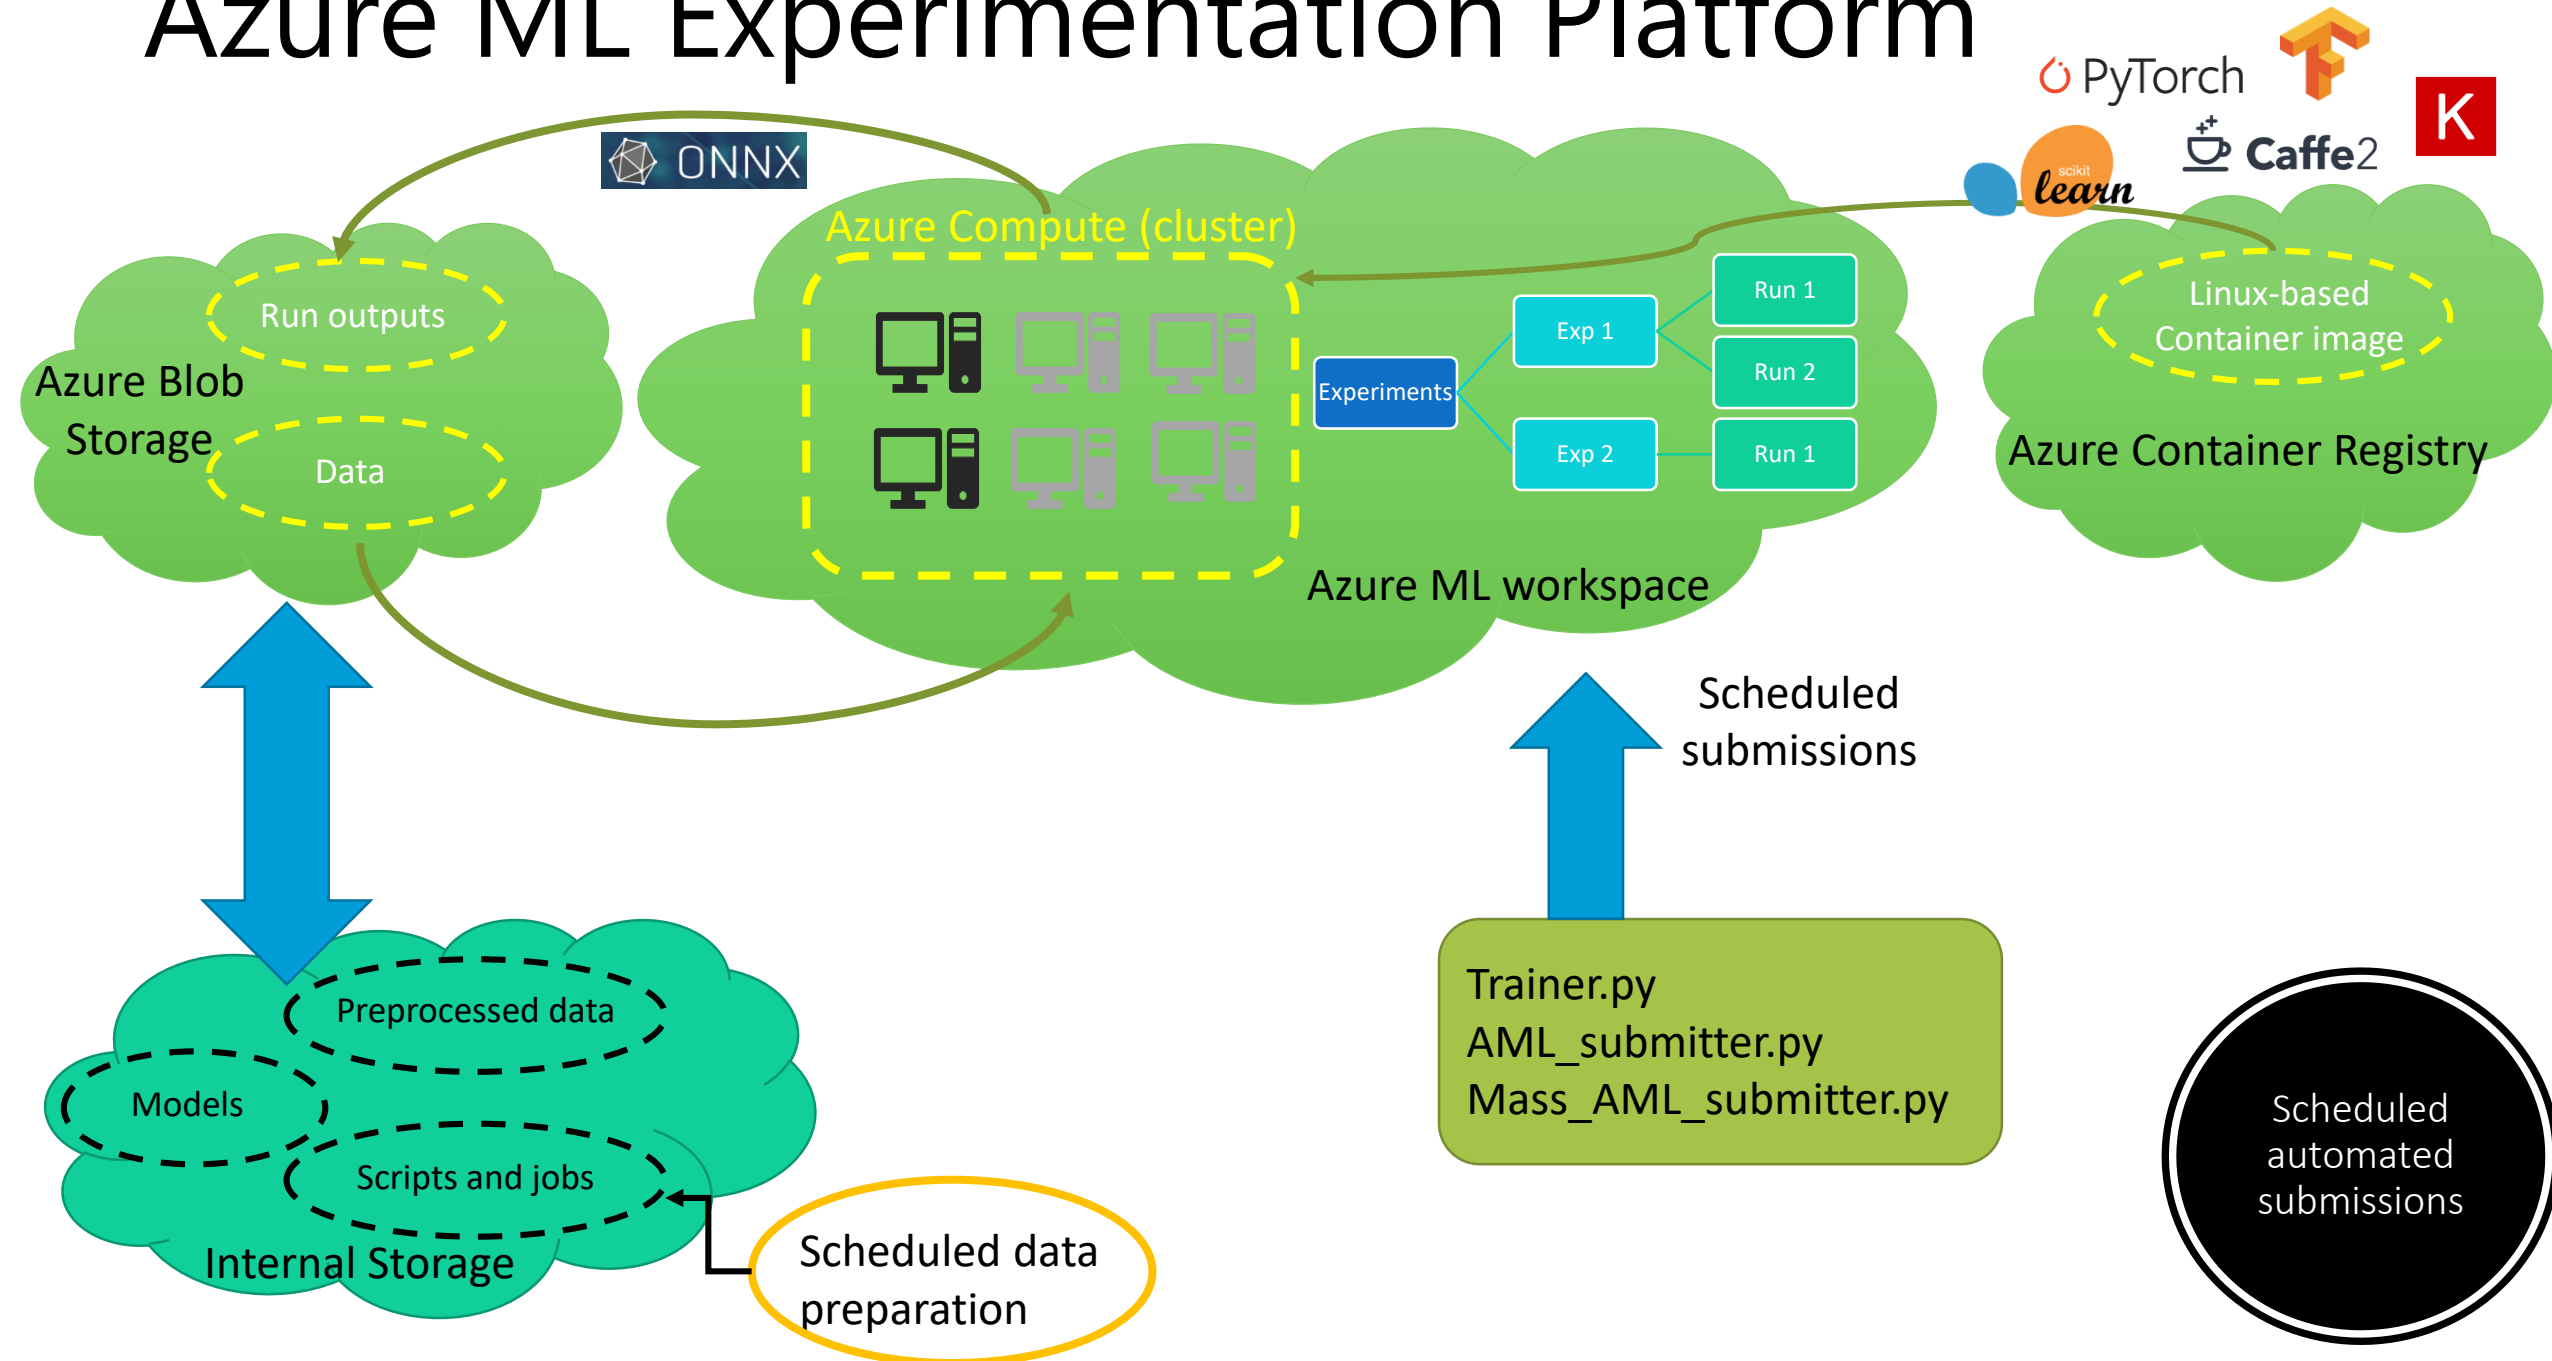
- Sensitive to time of day, day of week
- Custom self-learning model per machine
- Precision of 94%



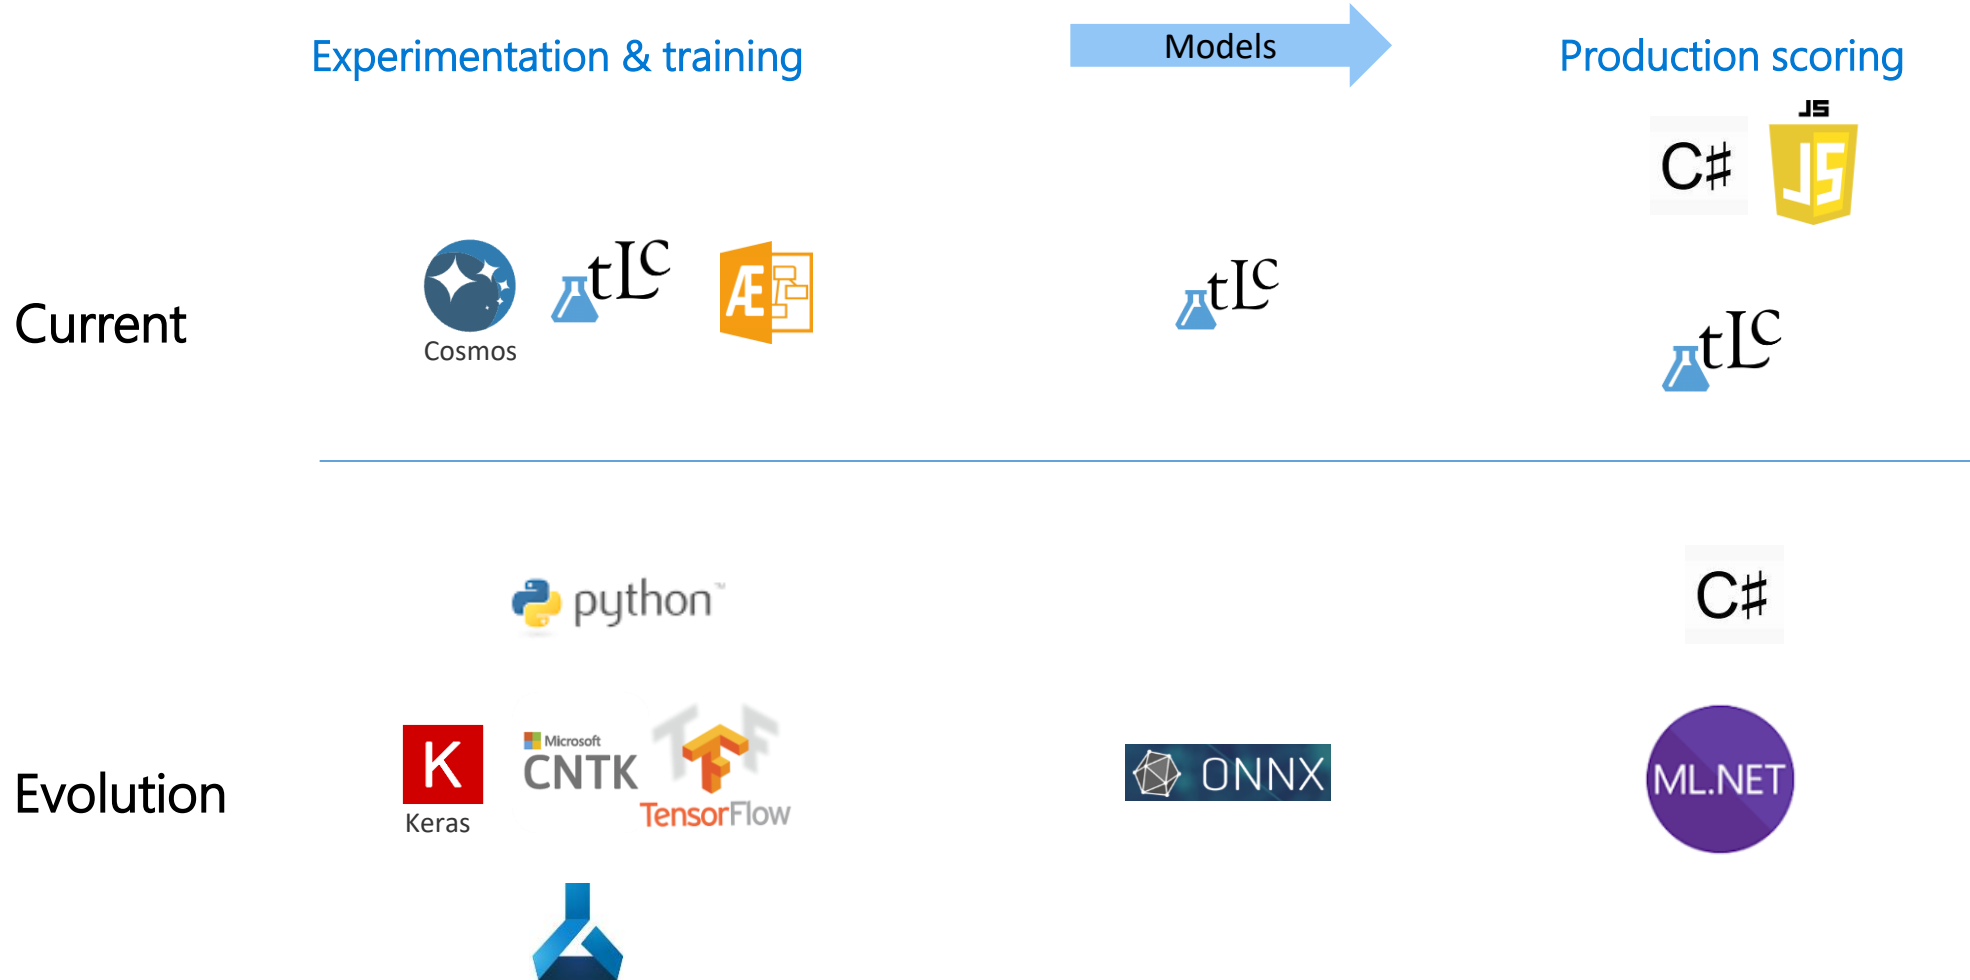


Tools and platform

Azure ML Experimentation Platform



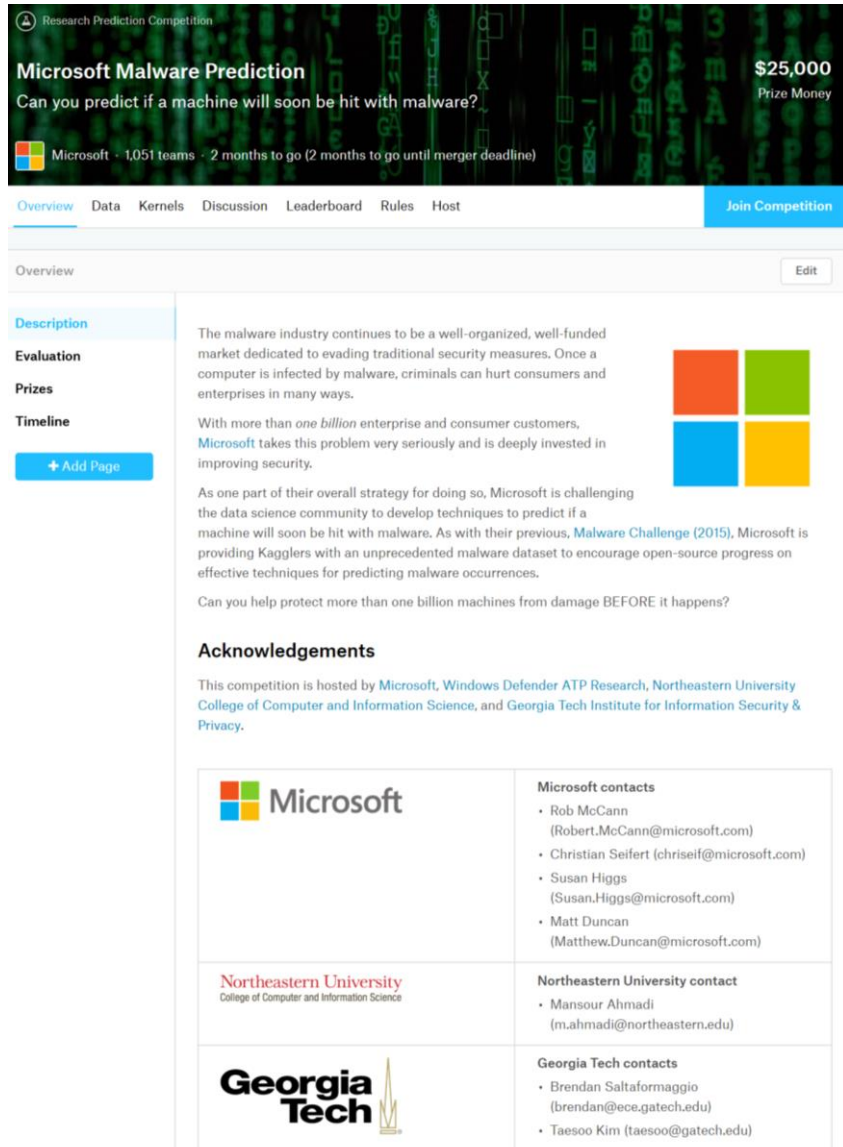
Tools of the trade





Community

Kaggle data science competition



The screenshot shows the 'Microsoft Malware Prediction' competition page on Kaggle. The header includes the competition title, a question 'Can you predict if a machine will soon be hit with malware?', a prize money of '\$25,000', and a progress bar showing 'Microsoft - 1,051 teams - 2 months to go (2 months to go until merger deadline)'. Navigation tabs include Overview, Data, Kernels, Discussion, Leaderboard, Rules, and Host. A 'Join Competition' button is visible. The 'Overview' section is expanded, showing a description of the malware industry, evaluation details, prizes, and a timeline. The 'Acknowledgements' section lists the hosts: Microsoft, Northeastern University, and Georgia Tech. A table at the bottom provides contact information for each organization.

| Organization | Contact Information |
|-------------------------|--|
| Microsoft | Microsoft contacts <ul style="list-style-type: none">Rob McCann (Robert.McCann@microsoft.com)Christian Seifert (chriseif@microsoft.com)Susan Higgs (Susan.Higgs@microsoft.com)Matt Duncan (Matthew.Duncan@microsoft.com) |
| Northeastern University | Northeastern University contact <ul style="list-style-type: none">Mansour Ahmadi (m.ahmadi@northeastern.edu) |
| Georgia Tech | Georgia Tech contacts <ul style="list-style-type: none">Brendan Saltaformaggio (brendan@ece.gatech.edu)Taesoo Kim (taesoo@gatech.edu) |

New 2018-19 Competition: Anticipate malware based on machine state

- Effort started w/ [internship](#), competition running 12/13/18 - 3/13/19
- <https://www.kaggle.com/c/microsoft-malware-prediction>

Collaborative!

- Academic partners (Northeastern, Georgia Tech, UW, UW Tacoma)
- Microsoft partners (ILDC, MSRA)
- >2,426 teams & >300 forum discussion threads, > 3,000 posts
- Winning submissions are being reviewed and hold promise for product impact
 - Durability over time is a focus

Newsworthy!

- [Our blog](#)
- Academic/MS partner ann.
- [ZDNet](#), [Tom's Hardware](#)
- [Bleeping Computer](#), [Neowin](#), ...

Internships

Undergraduate and graduate internships

- 12 weeks paid internship
- Access to real-world attack data
- Work on cool problems
- Often partners with Microsoft Research

Published papers


- Danny Hendler, Shay Kels, and Amir Rubin. "Detecting malicious PowerShell commands using deep neural networks." ACM, 2018.
- Jack W. Stokes, De Wang, Mady Marinescu, Marc Marino, Brian Bussone. "Attack and Defense of Dynamic Analysis-Based, Adversarial Neural Malware Detection Models." MILCOM, 2018.
- Yehonatan Cohen, Danny Hendler, and Amir Rubin. "Detection of malicious webmail attachments based on propagation patterns." Knowledge-Based Systems, 2018.
- Rakshit Agrawal, Jack W. Stokes, Mady Marinescu, Karthik Selvaraj. "Robust Neural Malware Detection Models for Emulation Sequence Learning." MILCOM, 2018.
- Md Amran Siddiqui, Jack W. Stokes, Christian Seifert, Evan Argyle, Robert McCann, Joshua Neil, Justin Carroll. "Detecting Cyber Attacks Using Anomaly Detection with Explanations and Expert Feedback." ICASSP, 2019.
- Rakshit Agrawal, Jack W. Stokes, Mady Marinescu, Karthik Selvaraj. "Neural Sequential Malware Detection with Parameters." ICASSP 2018
- Rakshit Agrawal, Jack Stokes, Karthik Selvaraj, Mady Marinescu. "Attention In Recurrent Neural Networks For Ransomware Detection." ICASSP 2019

Contact:

✉ chriseif@microsoft.com

 [Linkedin](#)

 [cseifert](#)



Questions